

# Lab2

鄭余玄

## 1 Introduction

In this lab, we are going to implement both *EEGNet* and *DeepConvNet*. The training and testing dataset is from the BCI Competition with 2 classes and 2 channels. I will explain my implementation in detail in Section 2. In addition, experiments setup and results are described in Section 3.

## 2 Method

First, my script will parse arguments of hyper-parameters, such as learning rate, activation function, etc. The activation functions are mapped in a `ModuleDict` for the simplicity of further experiments:

```
1 Activations = nn.ModuleDict([
2     ['relu', nn.ReLU()],
3     ['lrelu', nn.LeakyReLU()],
4     ['elu', nn.ELU()],
5 ])
```

Then, the model is loaded to specified device (default: `cuda`), and the BCI dataset is loaded by `DataLoader` of *pytorch*. During training, the optimizer must set to be zero for each epoch, and then calculate the loss of model. In addition, it is safe to use `torch.no_grad()` to prevent from updating the model when calculating the accuracy.

In addition, since the loss function is `nn.CrossEntropyLoss` which combines both `nn.LogSoftmax` and `nn.NLLLoss`, there is no need to apply `nn.Softmax` layer in the model. For prediction, the `outputs` are logits, and thus should apply a `softmax` layer.

```
1 for epoch in range(args.epochs):
2     for i, (inputs, labels) in enumerate(train_loader, start=1):
```

```

3     # zero the parameter gradients
4     optimizer.zero_grad()
5
6     # forward + backward + optimize
7     outputs = net(inputs)
8     loss = criterion(outputs, labels)
9     loss.backward()
10    optimizer.step()
11
12    # batch training accuracy
13    with torch.no_grad():
14        predicted = torch.argmax(softmax(outputs), dim=1)
15        total += labels.size(0)
16        correct += (predicted == labels).sum().item()

```

On the other hand, after the convolution in the forward phase, the hidden tensor is flattened by `x.view(-1, dim)`.

## 2.1 EEGNet

In *EEGNet*, the output shape of the first convolution layer (`firstConv`) is the same as input size; thus the padding size is adjusted according to the kernel size. In contrast to *tensorflow*, *pytorch* can not specified the padding size of a convolution layer to be **SAME**. Therefore, I wrote a `Conv2dSame` inheriting `Conv2d` such that the padding is automatically adjusted.

We know that the output layer of a convolution layer has size:

$$H_{\text{out}} = \left\lfloor \frac{H_{\text{in}} + 2 \times \text{padding}_H - K_H}{\text{stride}_H} + 1 \right\rfloor$$

where  $K_H$  is the “full” space of the kernel  $K_H = \text{dilation}_H \times (k_H - 1) + 1$ , and  $k_H$  is the input kernel shape. As a result, for a **SAME** padding layer ( $H_{\text{in}} = H_{\text{out}}$ ), the padding is calculated as the following:

$$2 \times \text{padding}_H = \left( \left\lfloor \frac{H_{\text{in}} + \text{stride}_H - 1}{\text{stride}_H} \right\rfloor - 1 \right) * \text{stride}_H + K_H - H_{\text{in}}$$

Note that if the padding is odd, I choose to pad right and bottom in my implementation.

## 2.2 DeepConvNet

In *DeepConvNet*, there are four repeated similar parts of convolution layers in the same pattern. Therefore, it could be simplify as in the form of `nn.ModuleList`:

```
1 channels = [25, 50, 100, 200]
2 self.convs = nn.ModuleList([
3     nn.Sequential(
4         nn.Conv2d(in_channels, out_channels, kernel_size=(1, 5)),
5         nn.BatchNorm2d(out_channels, **bn_args),
6         Activations[act],
7         nn.MaxPool2d(kernel_size=(1, 2), stride=(1, 2)),
8         nn.Dropout(p),
9     ) for in_channels, out_channels in zip(channels[:-1], channels[1:])
10 ])
```

## 2.3 Activation Function

These activation functions are used in the following experiments:

$$\text{ReLU}(x) = \max(x, 0) \quad (1)$$

$$\text{LeakyReLU}(x, \text{neg\_slope} = 0.01) = \begin{cases} x & \text{if } x \geq 0 \\ \text{neg\_slope} \times x & \text{otherwise} \end{cases} \quad (2)$$

$$\text{ELU}(x, \alpha = 1.) = \max(0, x) + \min(0, \alpha \times (e^x - 1)) \quad (3)$$

In Figure 1, it shows that the *ReLU* activation function has relatively “hard” clip off the input value, while the *LeakyReLU* preserve some information in the negative part. In comparison, *ELU* utilizes the tail of an exponential function as the negative part of its activation, and thus its derivative is more natural when  $x$  is close to 0. All of these activation function could solve the vanishing gradient problem from the `softmax` activation function.

$$\frac{d}{dx}\text{ReLU}(x) = \begin{cases} 1 & \text{if } x \geq 0 \\ 0 & \text{otherwise} \end{cases} \quad (4)$$

$$\frac{d}{dx}\text{LeakyReLU}(x, \text{neg\_slope} = 0.01) = \begin{cases} 1 & \text{if } x \geq 0 \\ \text{neg\_slope} & \text{otherwise} \end{cases} \quad (5)$$

$$\frac{d}{dx}\text{ELU}(x, \alpha = 1.) = \begin{cases} 1 & \text{if } x \geq 0 \\ \text{ELU}(x) + \alpha & \text{otherwise} \end{cases} \quad (6)$$




Figure 1: Activation functions

### 3 Result

	ELU	ReLU	LeakyReLU
EEGNet	87.59%	89.17%	88.06%
DeepConvNet	81.48%	81.67%	81.76%

Table 1: (Best) Test Accuracy

The configuration to reproduce the following experiments are:

- 1 `python3 lab2.py --net EEGNet -lr 0.001 --weight_decay 0.06`  
`↪ --kernel_size 55 --epochs 1200 --batch_size 128`  
`↪ --activation elu`
- 2 `python lab2.py --net DeepConvNet -lr 0.001 --weight_decay 0.06`  
`↪ --batch_size 128 --epochs 300 --activation elu`

From Table 1, we can conclude that the performance of these three activation functions are similar while **LeakyReLU** performs slightly better. On the other hand, the optimizer used in all experiments is **Adam** which performs the similar as **RMSprop**.

The main efforts on *EEGNet* is (1) the kernel size of the **firstConv** (2) the weight decay of optimizer (3) **BatchNorm2d** parameters. In comparison, the *DeepConvNet* is too easy to over-fit; therefore, it performs relatively worse.

First, in the original paper of *EEGNet*, the kernel size of `firstConv` is 64. However, in this dataset, I figure out that 53 and 55 performs the best while a too large kernel size could not learn well. On the other hand, at first during the training phase, I did not notice the `nn.CrossEntropyLoss` had included the `softmax` function, but I got a acceptable test accuracy. It provides a clue that the weights of the model is too large, and therefore needs a weight decay in the optimizer.

### 3.1 EEGNet

```

1 EEGNet(
2     (firstConv): Sequential(
3         (0): Conv2dSame(1, 16, kernel_size=(1, 55), stride=(1, 1),
4             ↪ padding=(None, None), bias=False)
5         (1): BatchNorm2d(16, eps=1e-05, momentum=None, affine=True,
6             ↪ track_running_stats=False)
7     )
8     (depthwiseConv): Sequential(
9         (0): Conv2d(16, 32, kernel_size=(2, 1), stride=(1, 1),
10            ↪ groups=16, bias=False)
11         (1): BatchNorm2d(32, eps=1e-05, momentum=None, affine=True,
12            ↪ track_running_stats=False)
13         (2): ReLU()
14         (3): AvgPool2d(kernel_size=(1, 4), stride=(1, 4),
15            ↪ padding=0)
16         (4): Dropout(p=0.25)
17     )
18     (separableConv): Sequential(
19         (0): Conv2dSame(32, 32, kernel_size=(1, 16), stride=(1, 1),
20            ↪ padding=(None, None), bias=False)
21         (1): BatchNorm2d(32, eps=1e-05, momentum=None, affine=True,
22            ↪ track_running_stats=False)
23         (2): ReLU()
24         (3): AvgPool2d(kernel_size=(1, 8), stride=(1, 8),
25            ↪ padding=0)
26         (4): Dropout(p=0.25)
27     )
28     (classify): Sequential(
29         (0): Linear(in_features=736, out_features=2, bias=True)
30     )
31 )

```

## 3.2 DeepConvNet

```
1 DeepConvNet(  
2     (conv0): Sequential(  
3         (0): Conv2d(1, 25, kernel_size=(1, 5), stride=(1, 1))  
4         (1): Conv2d(25, 25, kernel_size=(2, 1), stride=(1, 1))  
5         (2): BatchNorm2d(25, eps=1e-05, momentum=None, affine=True,  
6             ↪ track_running_stats=False)  
7         (3): ReLU()  
8         (4): MaxPool2d(kernel_size=(1, 2), stride=(1, 2),  
9             ↪ padding=0, dilation=1, ceil_mode=False)  
10        (5): Dropout(p=0.8)  
11    )  
12    (convs): ModuleList(  
13        (0): Sequential(  
14            (0): Conv2d(25, 50, kernel_size=(1, 5), stride=(1, 1))  
15            (1): BatchNorm2d(50, eps=1e-05, momentum=None,  
16                ↪ affine=True, track_running_stats=False)  
17            (2): ReLU()  
18            (3): MaxPool2d(kernel_size=(1, 2), stride=(1, 2),  
19                ↪ padding=0, dilation=1, ceil_mode=False)  
20            (4): Dropout(p=0.8)  
21        )  
22        (1): Sequential(  
23            (0): Conv2d(50, 100, kernel_size=(1, 5), stride=(1, 1))  
24            (1): BatchNorm2d(100, eps=1e-05, momentum=None,  
25                ↪ affine=True, track_running_stats=False)  
26            (2): ReLU()  
27            (3): MaxPool2d(kernel_size=(1, 2), stride=(1, 2),  
28                ↪ padding=0, dilation=1, ceil_mode=False)  
29            (4): Dropout(p=0.8)  
30        )  
31        (2): Sequential(  
32            (0): Conv2d(100, 200, kernel_size=(1, 5), stride=(1, 1))  
33            (1): BatchNorm2d(200, eps=1e-05, momentum=None,  
34                ↪ affine=True, track_running_stats=False)  
35            (2): ReLU()  
36            (3): MaxPool2d(kernel_size=(1, 2), stride=(1, 2),  
37                ↪ padding=0, dilation=1, ceil_mode=False)  
38            (4): Dropout(p=0.8)  
39        )  
40    )  
41 )
```

```

32     )
33     (classify): Sequential(
34         (0): Linear(in_features=8600, out_features=2, bias=True)
35     )
36 )

```


### 3.3 Comparison Figure

The comparison figure is shown in Figure 2. In both models, the train and test accuracy converges in few epochs. As mentioned above, the training accuracy of *DeepConvNet* converges to 100% very fast, that is overfitting, and the test accuracy is almost not increasing after 100 epochs in contrast to the *EEGNet* which could slightly increase its test accuracy even at 1200 epochs.

## 4 Discussion


In this lab, I have tried a lot of skills to improve the performance of models. The most significant result is the *EEGNet* which has a more shallow neural network in comparison to *DeepConvNet*. In contrast, all methods I have had applied on the *DeepConvNet* have little effects and all over-fit in a small number of epochs, even if the dropout ratio is set to 0.99. These could imply that a shallow neural network would perform better under a limited small size of dataset.

## EEGNet



(a) EEGNet Accuracy

## DeepConvNet



(b) DeepConvNet Accuracy

Figure 2: Accuracy